

```

%matplotlib inline

import datetime
import itertools
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import tensorflow as tf

# How to enable Colab GPUs https://colab.research.google.com/notebooks/gpu.ipynb
# Select the Runtime > "Change runtime type" menu to enable a GPU accelerator,
# and then re-execute this cell.
if 'google.colab' in str(get_ipython()):
    device_name = tf.test.gpu_device_name()
    if device_name != '/device:GPU:0':
        raise SystemError('GPU device not found')
    print('Found GPU at: {}'.format(device_name))

import tensorflow.keras as keras
from keras.models import Sequential
from keras.layers import Input, Dense, Dropout, Activation, Conv1D, \
    BatchNormalization, GlobalAveragePooling1D, Flatten, \
    Reshape, LSTM
from keras.optimizers import Adam, Adadelta
from keras.callbacks import EarlyStopping, ReduceLROnPlateau

# Reduces variance in results but won't eliminate it :-(
%env PYTHONHASHSEED=0
import random
random.seed(42)
np.random.seed(42)
tf.random.set_seed(42)

Found GPU at: /device:GPU:0
env: PYTHONHASHSEED=0

```

✓ Keras Neural Networks for Weather Time Series Nowcasts

Building neural networks with [keras](#) for time series analysis of Cambridge UK weather data, using a streamlined version of data preparation from [Tensorflow time series forecasting tutorial](#).

Import Data

Data has been cleaned but may still have issues. See the [cleaning section](#) in the [Cambridge Temperature Model](#) repository for details.

The `y` variable is temperature * 10. I'm primarily interested in short term temperature forecasts (less than 2 hours). Observations occur every 30 mins.

```

if 'google.colab' in str(get_ipython()):
    data_loc = "https://github.com/makeyourownmaker/CambridgeTemperatureModel/blob/master/data/CamUKWeather.csv?raw=true"
else:
    data_loc = "../data/CamUKWeather.csv"
df = pd.read_csv(data_loc, parse_dates = True)

df['ds'] = pd.to_datetime(df['ds'])

print("Shape:")
print(df.shape)
print("\nInfo:")
print(df.info())
print("\nSummary stats:")
display(df.describe())
print("\nRaw data:")
display(df)
print("\n")

def plot_examples(data, x_var):
    """Plot 9 sets of observations in 3 * 3 matrix ..."""

    assert len(data) == 9

    cols = [col for col in data[0].columns if col != x_var]

    fig, axs = plt.subplots(3, 3, figsize = (15, 10))

```

```

    axs = axs.ravel() # apl for the win :-)

    for i in range(9):
        for col in cols:
            axs[i].plot(data[i][x_var], data[i][col])
            axs[i].xaxis.set_tick_params(rotation = 20, labelsiz = 10)

    fig.legend(cols, loc = 'upper center', ncol = len(cols))

    return None

cols = ['ds', 'y', 'humidity', 'dew.point', 'pressure',
        'wind.speed.mean', 'wind.bearing.mean', 'wind.speed.max']
plots = 9
window = 24
starts = [random.randint(0, np.floor(df.shape[0] / window)) for _ in range(plots)]
p_data = [df.loc[starts[i] * window:starts[i] * window + window, cols]
           for i in range(plots)]
plot_examples(p_data, 'ds')

```

Shape:
(192885, 11)

Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 192885 entries, 0 to 192884
Data columns (total 11 columns):
Column Non-Null Count Dtype

0 ds 192885 non-null datetime64[ns]
1 year 192885 non-null int64
2 doy 192885 non-null int64
3 time 192885 non-null object
4 y 192885 non-null int64
5 humidity 192885 non-null int64
6 dew.point 192885 non-null int64
7 pressure 192885 non-null int64
8 wind.speed.mean 192885 non-null int64
9 wind.bearing.mean 192885 non-null int64
10 wind.speed.max 192885 non-null int64
dtypes: datetime64[ns](1), int64(9), object(1)
memory usage: 16.2+ MB
None

Summary stats:

	year	doy	y	humidity	dew.point	pressure	wind.speed.mean	wind.bearing.mean
count	192885.000000	192885.000000	192885.000000	192885.000000	192885.000000	192885.000000	192885.000000	192885.000000
mean	2013.895803	186.882298	101.096819	79.239951	62.135174	1014.404153	44.588148	196.223423
std	3.283992	106.486420	64.465602	16.908724	51.016879	11.823922	40.025546	82.458390
min	2008.000000	1.000000	-138.000000	25.000000	-143.000000	963.000000	0.000000	0.000000
25%	2011.000000	94.000000	52.000000	69.000000	25.000000	1008.000000	12.000000	135.000000
50%	2014.000000	191.000000	100.000000	83.000000	64.000000	1016.000000	35.000000	225.000000
75%	2017.000000	280.000000	145.000000	92.000000	100.000000	1023.000000	67.000000	270.000000
max	2020.000000	366.000000	361.000000	100.000000	216.000000	1048.000000	291.000000	315.000000

Raw data:

	ds	year	doy	time	y	humidity	dew.point	pressure	wind.speed.mean	wind.bearing.mean	wind.speed.max
0	2008-08-01 08:30:00	2008	214	09:30:00	186	69	128	1010	123	180	280
1	2008-08-01 09:00:00	2008	214	10:00:00	191	70	135	1010	137	180	260
2	2008-08-01 09:30:00	2008	214	10:30:00	195	68	134	1010	133	180	260
	2008-08-										

✓ Data Processing and Feature Engineering

The data must be reformatted before model building.

The following steps are necessary:

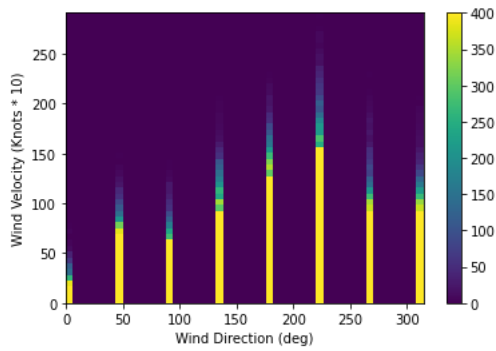
- Wind direction and speed transformation
- Impute missing data where possible
- Time conversion
- Split data
- Normalise data
- Window data

Wind direction and speed transformation

The `wind.bearing.mean` column gives wind direction in degrees but is categorised at 45 degree increments, i.e. 0, 45, 90, 135, 180, 225, 270, 315. Wind direction shouldn't matter if the wind is not blowing.

The distribution of wind direction and speed looks like this:

```
02:00:00  
plt.hist2d(df['wind.bearing.mean'], df['wind.speed.mean'], bins = (50, 50), vmax = 400)  
plt.colorbar()  
plt.xlabel('Wind Direction (deg)')  
plt.ylabel('Wind Velocity (Knots * 10)');
```



Convert wind direction and speed to x and y vectors, so the model can more easily interpret them.

```

wv = df['wind.speed.mean']
max_wv = df['wind.speed.max']

# Convert to radians
wd_rad = df['wind.bearing.mean'] * np.pi / 180


# Calculate the wind x and y components
df['wind.x'] = wv * np.cos(wd_rad)
df['wind.y'] = wv * np.sin(wd_rad)

# Calculate the max wind x and y components
df['max.wind.x'] = max_wv * np.cos(wd_rad)
df['max.wind.y'] = max_wv * np.sin(wd_rad)

df_orig = df

plt.hist2d(df['wind.x'], df['wind.y'], bins = (50, 50), vmax = 400)
plt.colorbar()
plt.xlabel('Wind X (Knots * 10)')
plt.ylabel('Wind Y (Knots * 10)')
plt.title('Wind velocity');

```



Better, but not ideal. Data augmentation with the [mixup](#) method is carried out at batch preparation time below.

From the paper: "mixup trains a neural network on convex combinations of pairs of examples and their labels".

Further details on how I apply mixup to time series are included in the Window data section below.

Here is an illustration of the improvement in wind velocity with mixup augmentation.

```

def mixup(data, alpha = 1.0, factor = 1):
    batch_size = len(data) - 1

    data['epoch'] = data.index.astype(np.int64) // 10**9

    # random sample lambda value from beta distribution
    l = np.random.beta(alpha, alpha, batch_size * factor)
    X_l = l.reshape(batch_size * factor, 1)

    # Get a pair of inputs and outputs
    y1 = data['y'].shift(-1).dropna()
    y1_ = pd.concat([y1] * factor)

    y2 = data['y'][0:batch_size]
    y2_ = pd.concat([y2] * factor)

    X1 = data.drop('y', 1).shift(-1).dropna()

```

```

X1_ = pd.concat([X1] * factor)

X2 = data.drop('y', 1)
X2 = X2[0:batch_size]
X2_ = pd.concat([X2] * factor)

# Perform mixup
X = X1_ * X_1 + X2_ * (1 - X_1)
y = y1_ * 1 + y2_ * (1 - 1)


df = pd.DataFrame(y).join(X)
df = data.append(df).sort_values('epoch', ascending = True)
df = df.drop('epoch', 1)

df = df.drop_duplicates(keep = False)

return df

df_mix = mixup(df.loc[:, ['y', 'wind.x', 'wind.y']], factor = 2)
plt.hist2d(df_mix['wind.x'], df_mix['wind.y'], bins = (50, 50), vmax = 400)
plt.colorbar()
plt.xlabel('Wind X (Knots * 10)')
plt.ylabel('Wind Y (Knots * 10)')
plt.title('Wind velocity with mixup data augmentation');

```



Mixup improves the categorical legacy of the wind velocity data. Unfortunately, if outliers are present their influence will be reinforced.

✓ Missing value interpolation

Missing data is much less of a problem for [prophet models](#) which handle it seamlessly.

Currently there are around 8,000 missing observations in approximately 600 sections or "gaps". The gaps range in length from 30 mins to 45 days.

Gaps have length across variables, i.e. if there is a gap of length 6 (3 hours) then all variables are missing for 6 consecutive observations.

Missing observations can be [imputed](#).

Here I use a variation on [cubic interpolation](#). Vanilla cubic interpolation overshoots the data which introduces outlier values. I limit the gap length to 12 observations (6 hours).

Longer gaps will be accounted for at the train, test, validation split stage.

```

del_cols = ['doy', 'wind.bearing.mean', 'wind.speed.mean', 'wind.speed.max']
df_ts = df_orig.set_index('ds', drop = False)
df_ts.drop(del_cols, axis = 1, inplace = True)

# Add NaN values for missing observations
df_ts_na = df_ts.asfreq('30min')

# Set NaN year, time values using index
df_ts_na.ds = df_ts_na.index
df_ts_na.year = df_ts_na.index.year
df_ts_na.time = df_ts_na.index.time

# Count number of consecutive missing values
# There are more elegant ways to do this but they don't cope well with NaNs
# As far as I can tell, neither numpy nor pandas have native run length encoding functions
len_holes = pd.Series([len(list(g)) for k, g in itertools.groupby(df_ts_na.y.isnull()) if k])
len_holes_long = pd.Series(list(itertools.repeat(1, 1)) for 1 in len_holes)
len_holes_flat = pd.Series(list(itertools.chain(*len_holes_long)))

df_ts_na['missing_len'] = -100
df_ts_na['missing_len'] = df_ts_na['missing_len'].astype('Int64')

```

```

df_ts_na.loc[df_ts_na.y.isnull(), 'missing_len'] = len_holes_flat.to_numpy()

# Mark 24 observations before and after each group of NaNs - for plotting
df_ts_na['around_nan'] = -100
df_ts_na['around_nan'] = df_ts_na['around_nan'].astype('Int64')

for i in range(-24, 25):
    df_ts_na.loc[df_ts_na.y.isna().shift(i).fillna(False), 'around_nan'] = i

df_ts_na.loc[df_ts_na.y.isna(), 'around_nan'] = 0

# Interpolate - method = 'spline' very slow :-(
#             cubic and quadratic overshoot the data and introduce outliers
limit = 12
method = 'pchip'
for v in ['y', 'humidity', 'dew.point', 'pressure', 'wind.x', 'wind.y', 'max.wind.x', 'max.wind.y']:
    df_ts_na[v] = df_ts_na[v].interpolate(method = method, limit = limit)

# Extract missing observations and surrounding values into dict of lists for checking & plotting
# slow :-(
j = miss_len = 0
inner_list = []
miss_plus = {}
for index, row in df_ts_na.iterrows():
    if (row['around_nan'] > -25) | (row['missing_len'] > 0):
        inner_list.append(index)
        if row['missing_len'] > 0:
            miss_len = row['missing_len']
        j = 1
    else:
        if j == 1:
            miss_plus.setdefault(miss_len, []).append(inner_list)
            inner_list = []
        j = miss_len = 0

# print("keys: ", len(miss_plus.keys()))
# print("sum: ", sum(miss_plus.keys()))
# print("keys: ", sorted(miss_plus.keys()))
# print(len(miss_plus[29]))
# print(len(miss_plus[29][0]))
# print(miss_plus[29][0])
# df_ts_na.loc[miss_plus[29][0]]

def plot_interpolations(data):
    """Plot 8 labeled interpolation examples in 3 x 3 subplots"""

    fig, axs = plt.subplots(3, 3, figsize = (15, 10), tight_layout = True)
    axs = axs.ravel()
    i = 0

    for v in ['y', 'humidity', 'dew.point', 'pressure', 'wind.x', 'wind.y', 'max.wind.x', 'max.wind.y']:
        marks = data.loc[data.missing_len > 0, v]
        title = str(len(marks)) + ' ' + v + ' Interpolation(s)'

        axs[i].plot(data[v], marker = '.', linestyle = '--')
        axs[i].plot(marks, marker = '.', linestyle = '--')
        axs[i].xaxis.set_tick_params(rotation = 30, labelsz = 10)
        axs[i].set_title(title)
        axs[i].set_ylabel(v)
        axs[i].set_xlabel('Time')
        i += 1

    axs[i].set_visible(False)

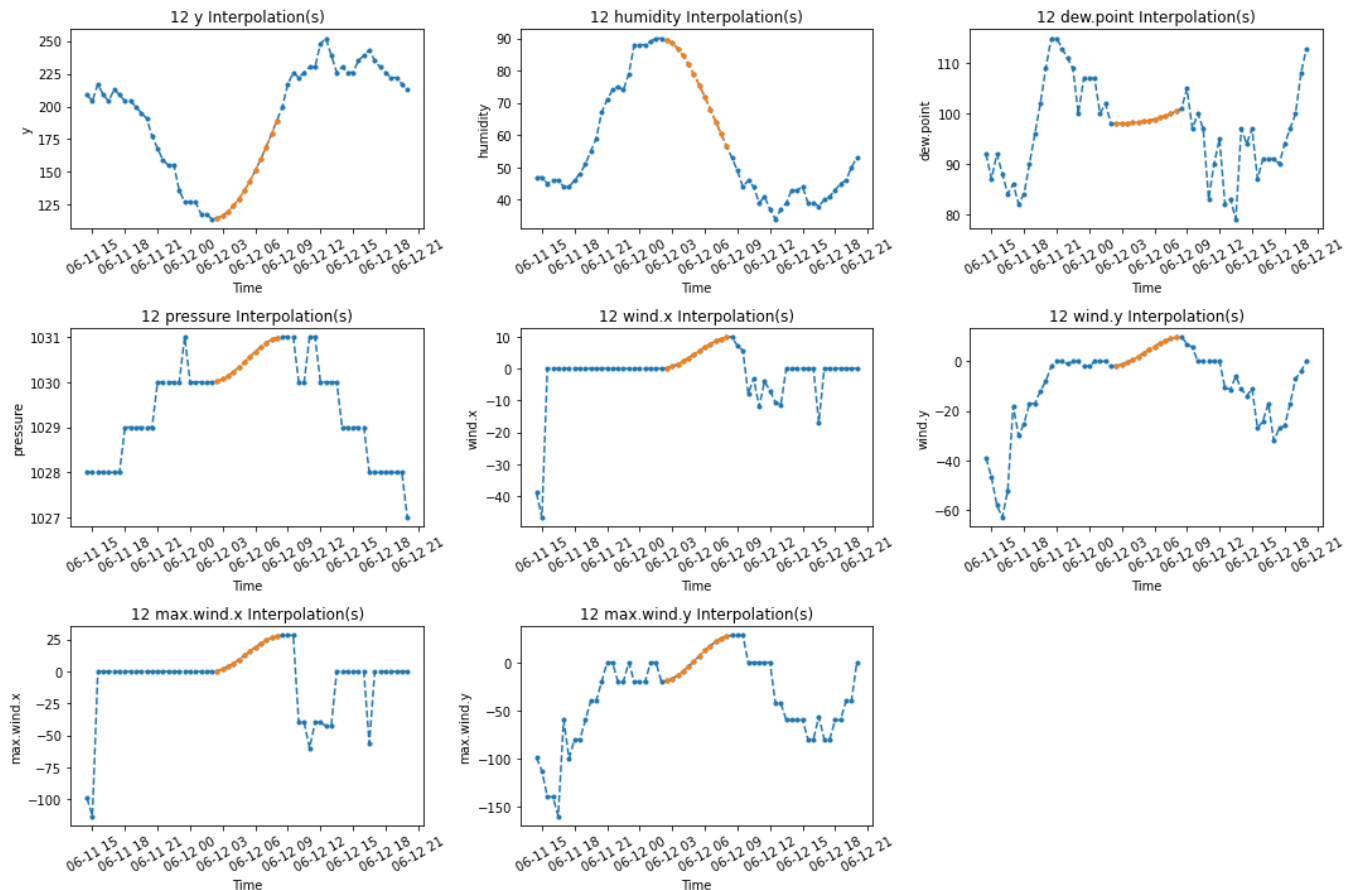
    return None

plot_interpolations(df_ts_na.loc[miss_plus[12][1]])

# Remove gaps longer than 12
df_ts_na = df_ts_na.dropna()
drop_cols = ['missing_len', 'around_nan']
df_ts_na.drop(drop_cols, axis = 1, inplace = True)
df = df_ts_na

```

```
/usr/local/lib/python3.7/dist-packages/scipy/interpolate/_cubic.py:293: RuntimeWarning: invalid value encountered in add
whmean = (w1/mk[:-1] + w2/mk[1:]) / (w1 + w2)
```



There is a 23 % reduction in missing values.

Alternative interpolation methods (piecewise methods in particular) may give more natural results for the wind and pressure variables.

✓ Time conversion


Convert `ds` timestamps to "time of day" and "time of year" variables using `sin` and `cos`.

```
# Convert to secs
date_time = pd.to_datetime(df['ds'], format = '%Y.%m.%d %H:%M:%S')
timestamp_s = date_time.map(datetime.datetime.timestamp)

day = 24 * 60 * 60
year = (365.2425) * day

df['day.sin'] = np.sin(timestamp_s * (2 * np.pi / day))
df['day.cos'] = np.cos(timestamp_s * (2 * np.pi / day))
df['year.sin'] = np.sin(timestamp_s * (2 * np.pi / year))
df['year.cos'] = np.cos(timestamp_s * (2 * np.pi / year))

plt.plot(np.array(df['day.sin'])[0:49])
plt.plot(np.array(df['day.cos'])[0:49])
plt.xlabel('Time')
plt.legend(['sin', 'cos'], loc = 'lower right')
plt.title('Time of day signal');
```



✓ Split data

Use data from 2018 for validation and 2019 for testing. These are entirely arbitrary choices. This results in an approximate 82%, 9%, 9% split for the training, validation, and test sets.

```
keep_cols = ['y', 'humidity', 'dew.point', 'pressure', 'wind.x', 'wind.y',
             'day.sin', 'day.cos', 'year.sin', 'year.cos']
del_cols = ['ds', 'time', 'max.wind.x', 'max.wind.y']
df.drop(del_cols, axis = 1, inplace = True)

train_df = df.loc[(df['year'] != 2018) & (df['year'] != 2019)]
valid_df = df.loc[df['year'] == 2018]
test_df = df.loc[df['year'] == 2019]

train_df = train_df.drop('year', axis = 1) # inplace = True gives SettingWithCopyWarning
valid_df = valid_df.drop('year', axis = 1) # ...
test_df = test_df.drop('year', axis = 1)
df = df.drop('year', axis = 1)

print("df.drop shape: ", df.shape)
print("train shape: ", train_df.shape)
print("valid shape: ", valid_df.shape)
print("test shape: ", test_df.shape)

df.drop shape: (194736, 10)
train shape: (160059, 10)
valid shape: (17236, 10)
test shape: (17441, 10)
```

✓ Normalise data

Features should be scaled before neural network training. Arguably, scaling should be done using moving averages to avoid accessing future values.

Instead, simple [standard score](#) normalisation will be used.

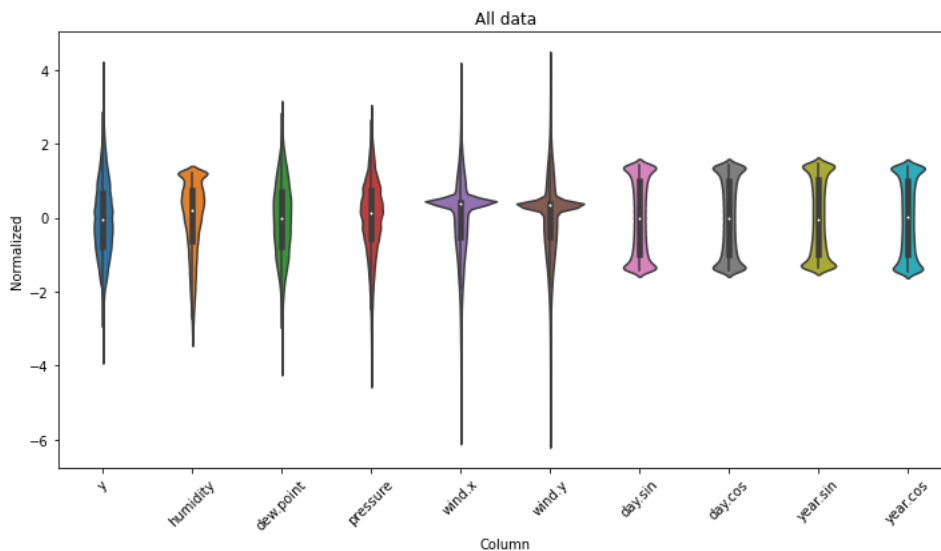
Plot [violin plot](#) to see distribution of features.

```
train_mean = train_df.mean()
train_std = train_df.std()

train_df = (train_df - train_mean) / train_std
valid_df = (valid_df - train_mean) / train_std
test_df = (test_df - train_mean) / train_std

df_std = (df - train_mean) / train_std
df_std = df_std.melt(var_name = 'Column', value_name = 'Normalized')

plt.figure(figsize=(12, 6))
ax = sns.violinplot(x = 'Column', y = 'Normalized', data = df_std)
ax.set_xticklabels(df.keys(), rotation = 45)
ax.set_title('All data');
```

There may still be some outliers present but there are no glaring problems.

✓ Window data

Models are trained using sliding windows of samples from the data.

Window parameters to consider for the [tf.keras.preprocessing.timeseries_dataset_from_array](#) function:

- `sequence_length`:
 - Length of the output sequences (in number of timesteps), or number of **lag** observations to use
- `sequence_stride`:
 - Period between successive output sequences. For stride `s`, output samples start at index `data[i]`, `data[i + s]`, `data[i + 2 * s]` etc
 - `s` can include an **offset** and/or 1 or more **steps ahead** to forecast
- `batch_size`:
 - Number of samples in each batch
- `shuffle`:
 - Shuffle output samples, or use chronological order

Initial values used:

- `sequence_length` (aka lags): 24 (corresponds to 12 hours)
- `steps ahead` (what to forecast): 1 and separately 4 (corresponds to 30 mins and separately 30 mins, 60 mins, 90 mins, 120 mins)
- `offset` (space between lags and steps ahead): 0
- `batch_size`: 32
- `shuffle`: True for training data

The `make_dataset` function below generates [tensorflow datasets](#) for:

- Lags, steps-ahead, offset, batch size and shuffle
- Optionally multiple y columns (Not extensively tested)

Stride is used to specify offset + steps-ahead. Offset will be 0 throughout this notebook.

TODO Insert figure illustrating lags, offsets and steps-ahead.

`shuffle = True` is used with train data. `shuffle = False` is used with validation and test data so the residuals can be checked for heteroscedasticity.

Throughout this notebook I use a shorthand notation to describe lags and strides. For example:

- 24l_1s is 24 lags 1 step ahead
- 24l_4s is 24 lags 4 steps ahead

Mixup data augmentation

Data augmentation with [mixup: Beyond Empirical Risk Minimization](#) by Zhang *et al* is used to help counter the categorical legacy from the wind bearing variable. Simple 'input mixup' is used as opposed to the batch-based mixup Zhang *et al* focus on. Input mixup has the advantage that it can be used with non-neural network methods. Mixup is performed for train and validation data separately. With current settings these datasets are approximately 3 times larger but this can be varied. Three times more training data is manageable on Colab. Test and validation data is left unmodified.

I apply mixup between consecutive observations in the time series instead of the usual random observations. This is a fairly conservative starting point. I'd be surprised if applying mixup between consecutive days of measurements didn't give better results. Applying mixup between inputs with equal temperature values will not improve performance and will increase run time.

TODO Insert couple of examples of mixup - use `plot_examples()`

I don't show it in this notebook, but adding this data augmentation makes a big difference to loss values for all three model architectures. For example, here are comparable results for MLP, 24 lags, 1 step ahead, 20 epochs.

Augmentation	Train rmse	Train mae	Valid rmse	Valid mae
No augmentation	0.0058	0.053	0.0054	0.052
Input mixup	0.0016	0.025	0.0015	0.025

See this [commit](#) for results from other architectures without input mixup.

```
def make_dataset(data, y_cols, lags = 1, steps_ahead = 1, stride = 1, bs = 32, shuffle = False, mix = 2):
    assert stride >= steps_ahead

    total_window_size = lags + stride

    # Add NaN values for missing observations
    data = data.asfreq('30min')

    # Split data into subsets (blocks) on NaNs - SLOW - https://stackoverflow.com/a/21404655/100129
    blocks = np.split(data, np.where(np.isnan(data.y))[0])
    # Removing NaN entries
    blocks = [bl[~np.isnan(bl.y)] for bl in blocks if not isinstance(bl, np.ndarray)]
    # Removing empty DataFrames
    blocks = [bl for bl in blocks if not bl.empty]

    i = 0
    for block in blocks:
        i += 1
        if mix != 0:
            block_mix = mixup(block, factor = mix)
        else:
            block_mix = block
        block_np = np.array(block_mix, dtype = np.float32)

        ds = tf.keras.preprocessing.timeseries_dataset_from_array(
            data = block_np,
            targets = None,
            sequence_length = total_window_size,
            sequence_stride = 1,
            shuffle = shuffle,
            batch_size = bs)

        col_indices = {name: i for i, name in enumerate(data.columns)}
        X_slice = slice(0, lags)
        y_start = total_window_size - steps_ahead
        y_slice = slice(y_start, None)

    def split_window(features):
        X = features[:, X_slice, :]
        y = features[:, y_slice, :]

        X = tf.stack(
            [X[:, :, col_indices[name]] for name in data.columns],
            axis = -1)
        y = tf.stack(
            [y[:, :, col_indices[name]] for name in y_cols],
            axis = -1)

        # Slicing doesn't preserve static shape information, so set the shapes manually.
        # This way the `tf.data.Datasets` are easier to inspect.
        X.set_shape([None, lags, None])
        y.set_shape([None, steps_ahead, None])

        return X, y

    ds = ds.map(split_window)

    if i == 1:
        combined_dataset = ds
    else:
        combined_dataset = combined_dataset.concatenate(ds)

    return combined_dataset
```

```

def make_datasets(train, valid, test,
                 y_cols = 'y', lags = 1, steps_ahead = 1,
                 stride = 1, bs = 32, shuffle = False):
    ds_train = make_dataset(train, y_cols,
                           lags = lags, steps_ahead = steps_ahead,
                           stride = stride, shuffle = shuffle)
    ds_valid = make_dataset(valid, y_cols,
                           lags = lags, steps_ahead = steps_ahead,
                           stride = stride, shuffle = False)
    ds_test = make_dataset(test, y_cols,
                           lags = lags, steps_ahead = steps_ahead,
                           stride = stride, shuffle = False, mix = 0)

    return ds_train, ds_valid, ds_test

def dataset_sanity_checks(data, name):
    print(name, "batches: ", data.cardinality().numpy())
    for batch in data.take(1):
        print("\tx (batch_size, time, features): ", batch[0].shape)
        print("\ty (batch_size, time, features): ", batch[1].shape)
        print("\tx[0][0]: ", batch[0][0])
        print("\ty[0][0]: ", batch[1][0])

# Single step-ahead
ds = {}
bs = 32
shuffle = True
ds['train_24l_1s'], ds['valid_24l_1s'], ds['test_24l_1s'] = make_datasets(train_df,
                                                                           valid_df,
                                                                           test_df,
                                                                           lags = 24,
                                                                           shuffle = shuffle,
                                                                           bs = bs)

# dataset_sanity_checks(ds_train_4l_1s, '4l 1s train');

# 4 steps-ahead
steps = stride = 4
ds['train_24l_4s'], ds['valid_24l_4s'], ds['test_24l_4s'] = make_datasets(train_df,
                                                                           valid_df,
                                                                           test_df,
                                                                           lags = 24,
                                                                           steps_ahead = steps,
                                                                           stride = stride,
                                                                           shuffle = shuffle,
                                                                           bs = bs)

# lags = 4
# display(train_df.head(lags + steps))
# dataset_sanity_checks(ds_train_4l_4s, '4l 4s train');


# Plot 9 examples from ds['train_24l_1s']
def plot_dataset_examples(dataset):
    fig, axs = plt.subplots(3, 3, figsize = (15, 10))
    axs = axs.ravel()

    for batch in dataset.take(1):
        for i in range(9):
            x = batch[0][i].numpy()
            axs[i].plot(x)

    fig.legend(range(1, 11), loc = 'upper center', ncol = 10);

plot_dataset_examples(ds['train_24l_1s'])

```



✓ Model Building

Model architectures considered:

- MLP
- [FCN](#)
- ResNet
- LSTM

The architectures considered were mostly inspired by those used by Wang *et al* in [Time Series Classification from Scratch with Deep Neural Networks: A Strong Baseline](#). They did not consider LSTMs. Initial hyperparameter settings came from [Deep learning for time series classification: a review](#).

I'm primarily interested in "now-casting" or forecasting in the next 1 or 2 hours. The following model outputs are investigated:

- Single step ahead - 30 mins
- Multi-step ahead - 30, 60, 90 and 120 mins

The training and validation code are stored in the `compile_fit_validate` function below.

Multi-layer perceptron

It is useful to check the performance of the multi-layer perceptron (MLP) before using more sophisticated models. The MLP is described in the `build_mlp_model` function below. It deviates from the Wang *et al*/Fawaz *et al* model. Specifically, I use a `Flatten` layer for the first layer to train on multiple input lags, reduce the number of layers from 3 to 2 and reduce the number of neurons in each layer from 500 to 64.

First, check single step-ahead predictions.

```
def compile_fit_validate(model, train, valid, optimizer, epochs = 5, verbose = 2):
    # Reduces variance in results but won't eliminate it :- (
    random.seed(42)
    np.random.seed(42)
    tf.random.set_seed(42)

    if optimizer.lower() == 'adadelata':
        opt = Adadelata(lr = 1.0)
    else:
        opt = Adam(lr = 0.001)

    es = EarlyStopping(monitor = 'val_loss', mode = 'min', verbose = 1, patience = 10)
    lr = ReduceLROnPlateau(monitor = 'val_loss', factor = 0.2, patience = 5, min_lr = 0.0001)
```

```

model.compile(optimizer = opt, loss = 'mse', metrics = ['mae', 'mape'])
h = model.fit(train, validation_data = valid,
              epochs = epochs, verbose = verbose, callbacks = [es, lr])

return h

def plot_history(h, name, epochs = 10):
    fig, axs = plt.subplots(1, 2, figsize = (9, 6), tight_layout = True)
    axs = axs.ravel()

    axs[0].plot(h.history['loss'])
    axs[0].plot(h.history['val_loss'])
    axs[0].set_title(name + ' loss')
    axs[0].set_xticklabels(range(1, epochs + 1))
    axs[0].set_xticks(range(0, epochs))
    axs[0].set_ylabel('loss')
    axs[0].set_xlabel('epoch')
    axs[0].legend(['train', 'valid'], loc = 'upper right')

    axs[1].plot(h.history['mape'])
    axs[1].plot(h.history['val_mape'])
    axs[1].set_title(name + ' mape')
    axs[1].set_xticks(range(0, epochs))
    axs[1].set_xticklabels(range(1, epochs + 1))
    axs[1].set_title(name + ' mape')
    axs[1].set_ylabel('mape')
    axs[1].set_xlabel('epoch')
    axs[1].legend(['train', 'valid'], loc = 'upper right')
    plt.show()

    return None

def print_min_loss(h, name):
    argmin_loss = np.argmin(np.array(h.history['loss']))
    argmin_val_loss = np.argmin(np.array(h.history['val_loss']))
    min_loss = h.history['loss'][argmin_loss]
    min_val_loss = h.history['val_loss'][argmin_val_loss]
    mape = h.history['mape'][argmin_loss]
    val_mape = h.history['val_mape'][argmin_val_loss]
    mae = h.history['mae'][argmin_loss]
    val_mae = h.history['val_mae'][argmin_val_loss]

    txt = "{0:s} {1:s} min loss: {2:f}\\tmae: {3:f}\\tmape: {4:f}\\tepoche: {5:d}"
    print(txt.format(name, "train", min_loss, mae, mape, argmin_loss + 1))
    print(txt.format(name, "valid", min_val_loss, val_mae, val_mape, argmin_val_loss + 1))
    print()

    return None

def get_io_shapes(data):
    for batch in data.take(1):
        in_shape = batch[0][0].shape
        out_shape = batch[1][0].shape

    return in_shape, out_shape

def build_mlp_model(name, data, neurons = 64):
    in_shape, out_shape = get_io_shapes(data)
    out_steps = out_shape[0]

    mlp = Sequential(name = name)

    mlp.add(Input(shape = in_shape))
    mlp.add(Flatten()) # Shape: (time, features) => (time*features)
    # mlp.add(Dropout(0.1))

    mlp.add(Dense(neurons, activation = 'relu'))
    # mlp.add(Dropout(0.1))

    mlp.add(Dense(neurons, activation = 'relu'))
    # mlp.add(Dropout(0.1))

    mlp.add(Dense(out_steps))
    mlp.add(Reshape([1, -1]))

    return mlp

```

```
def run_model(model, train, valid, optimizer = 'adam', epochs = 5):
    in_shape, out_shape = get_io_shapes(train)
    model_id = model.name + ' model - ' + str(in_shape[0]) + \
        ' lags ' + str(out_shape[0]) + ' steps-ahead - '

    model.summary()
    h = compile_fit_validate(model, train, valid, optimizer, epochs)
    plot_history(h, model_id, epochs)
    print_min_loss(h, model_id)

    return h

h = {} # history
name = 'MLP'
models = {}
models['mlp_24l_1s'] = build_mlp_model(name, ds['train_24l_1s'])
models['mlp_24l_4s'] = build_mlp_model(name, ds['train_24l_4s'])
```

✓ Learning rate finder

Leslie Smith was one of the first people to work on finding optimal learning rates for deep learning networks in [Cyclical Learning Rates for Training Neural Networks](#). Jeremy Howard from [fast.ai](#) popularised the learning rate finder used here.

Before building any models, I use a modified version of [Pavel Surmenok's Keras learning rate finder](#) to get reasonably close to the optimal learning rate. It's a single small class which I add support for tensorflow datasets to, customise the graphics and add a simple summary function to.

```
from keras.callbacks import LambdaCallback
import keras.backend as K
import math

class LRFinder:
    """
    Plots the change of the loss function of a Keras model when the learning rate is exponentially increasing.
    See for details:
    https://towardsdatascience.com/estimating-optimal-learning-rate-for-a-deep-neural-network-ce32f2556ce0
    """

    def __init__(self, model):
        self.model = model
        self.losses = []
        self.lrs = []
        self.best_loss = 1e9

    def on_batch_end(self, batch, logs):
        # Log the learning rate
        lr = K.get_value(self.model.optimizer.lr)
        self.lrs.append(lr)

        # Log the loss
        loss = logs['loss']
        self.losses.append(loss)

        # Check whether the loss got too large or NaN
        if batch > 5 and (math.isnan(loss) or loss > self.best_loss * 4):
            self.model.stop_training = True
            return

        if loss < self.best_loss:
            self.best_loss = loss

        # Increase the learning rate for the next batch
        lr *= self.lr_mult
        K.set_value(self.model.optimizer.lr, lr)

    def find_ds(self, train_ds, start_lr, end_lr, batch_size=64, epochs=1, **kw_fit):
        # If x_train contains data for multiple inputs, use length of the first input.
        # Assumption: the first element in the list is single input; NOT a list of inputs.
        # N = x_train[0].shape[0] if isinstance(x_train, list) else x_train.shape[0]
        N = train_ds.cardinality().numpy()

        # Compute number of batches and LR multiplier
        num_batches = epochs * N / batch_size
        self.lr_mult = (float(end_lr) / float(start_lr)) ** (float(1) / float(num_batches))
        #print(self.lr_mult)
        # Save weights into a file
```

```

initial_weights = self.model.get_weights()

# Remember the original learning rate
original_lr = K.get_value(self.model.optimizer.lr)

# Set the initial learning rate
K.set_value(self.model.optimizer.lr, start_lr)

callback = LambdaCallback(on_batch_end=lambda batch, logs: self.on_batch_end(batch, logs))

self.model.fit(train_ds,
                batch_size=batch_size, epochs=epochs,
                callbacks=[callback],
                **kw_fit)

# Restore the weights to the state before model fitting
self.model.set_weights(initial_weights)

# Restore the original learning rate
K.set_value(self.model.optimizer.lr, original_lr)

def plot_loss(self, axs, sma, n_skip_beginning=10, n_skip_end=5, x_scale='log'):
    """
    Plots the loss.
    Parameters:
        n_skip_beginning - number of batches to skip on the left.
        n_skip_end - number of batches to skip on the right.
    """
    lrs = self.lrs[n_skip_beginning:-n_skip_end]
    losses = self.losses[n_skip_beginning:-n_skip_end]
    best_lr = self.get_best_lr(sma=sma, n_skip_beginning=10, n_skip_end=5)

    axs[0].set_ylabel("loss")
    axs[0].set_xlabel("learning rate (log scale)")
    axs[0].plot(lrs, losses)
    axs[0].vlines(best_lr, np.min(losses), np.max(losses), linestyle='dashed')
    axs[0].set_xscale(x_scale)

def plot_loss_change(self, axs, sma=1, n_skip_beginning=10, n_skip_end=5, y_lim=None):
    """
    Plots rate of change of the loss function.
    Parameters:
        axs - subplot axes
        sma - number of batches for simple moving average to smooth out the curve.
        n_skip_beginning - number of batches to skip on the left.
        n_skip_end - number of batches to skip on the right.
        y_lim - limits for the y axis.
    """
    derivatives = self.get_derivatives(sma)[n_skip_beginning:-n_skip_end]
    lrs = self.lrs[n_skip_beginning:-n_skip_end]
    best_lr = self.get_best_lr(sma=sma, n_skip_beginning=n_skip_beginning, n_skip_end=n_skip_end)
    y_min, y_max = np.min(derivatives), np.max(derivatives)
    x_min, x_max = np.min(lrs), np.max(lrs)

    axs[1].set_ylabel("rate of loss change")
    axs[1].set_xlabel("learning rate (log scale)")
    axs[1].plot(lrs, derivatives)
    axs[1].vlines(best_lr, y_min, y_max, linestyle='dashed')
    axs[1].hlines(0, x_min, x_max, linestyle='dashed')
    axs[1].set_xscale('log')
    if y_lim == None:
        axs[1].set_ylim([y_min, y_max])
    else:
        axs[1].set_ylim(y_lim)

def get_derivatives(self, sma):
    assert sma >= 1
    derivatives = [0] * sma
    for i in range(sma, len(self.lrs)):
        derivatives.append((self.losses[i] - self.losses[i - sma]) / sma)
    return derivatives

def get_best_lr(self, sma, n_skip_beginning=10, n_skip_end=5):
    derivatives = self.get_derivatives(sma)
    best_der_idx = np.argmin(derivatives[n_skip_beginning:-n_skip_end])
    return self.lrs[n_skip_beginning:-n_skip_end][best_der_idx]

def summarise_lr(self, train_ds, start_lr, end_lr, batch_size=64, epochs=1, sma=1, n_skip_beginning=200, **kw_fit):
    self.find_ds(train_ds, start_lr, end_lr, batch_size, epochs)

    fig, axs = plt.subplots(1, 2, figsize = (9, 6), tight_layout = True)
    axs = axs.ravel()

```

```

self.plot_loss(axes, sma)
self.plot_loss_change(axes, sma=sma, n_skip_beginning=n_skip_beginning, n_skip_end=5)
plt.show()

```

```

best_lr = self.get_best_lr(sma=sma, n_skip_beginning=n_skip_beginning, n_skip_end=5)
print("best lr:", best_lr, "\n")

```


```

lrf = {}
model = models['mlp_24l_1s']
model.compile(optimizer = 'adadelat', loss = 'mse', metrics = ['mae', 'mape'])
lrf_mlp_24l_1s = LRFinder(model)
lrf_mlp_24l_1s.summarise_lr(ds['train_24l_1s'], 0.01, 1, 32, 5, 250, 25)
lrf['mlp_24l_1s'] = lrf_mlp_24l_1s

model = models['mlp_24l_4s']
model.compile(optimizer = 'adadelat', loss = 'mse', metrics = ['mae', 'mape'])
lrf_mlp_24l_4s = LRFinder(model)
lrf_mlp_24l_4s.summarise_lr(ds['train_24l_4s'], 0.01, 1, 32, 5, 250, 25)
lrf['mlp_24l_4s'] = lrf_mlp_24l_4s


```

Epoch 1/5
18904/18904 [=====] - 32s 2ms/step - loss: 13.8065 - mae: 0.2728 - mape: 118.5735



best lr: 0.014789497

Epoch 1/5
18869/18869 [=====] - 33s 2ms/step - loss: 21.2164 - mae: 0.3185 - mape: 152.5291



best lr: 0.014892996

The learning rate finder has a surprisingly low run time; possibly because the loss quickly becomes infinite at high learning rates.

The learning rate finder has not been very useful with any of these architectures and this data (see below for results from the other architectures). The models currently converge to the minimum loss value within 20 epochs with default learning rates. So, I default back to the accepted learning rate of 1.0 for adadelat and 0.001 for adam.

The smoothing value `sma`, is relatively high for the MLPs. It's possible to get lower rate of loss change values by using a lower `start_lr` but the rate of loss change has high variance in these regions. For MLPs learning rates in the region 0.01 to 1.0 give acceptable rates of loss change.

I leave the learning rate finder code in this notebook for possible future personal reference. It may also prove useful with other architectures and/or addition of exogenous regressors from for example the [Global Forecast System](#) model.

First, check single step-ahead model.

```
h['mlp_24l_1s'] = run_model(models['mlp_24l_1s'], ds['train_24l_1s'], ds['valid_24l_1s'], optimizer = 'adadelata', epochs = 20)
```

```
Model: "MLP"
```

Layer (type)	Output Shape	Param #
flatten_2 (Flatten)	(None, 240)	0
dense_6 (Dense)	(None, 64)	15424
dense_7 (Dense)	(None, 64)	4160
dense_8 (Dense)	(None, 1)	65
reshape_2 (Reshape)	(None, 1, 1)	0

```
Total params: 19,649
Trainable params: 19,649
Non-trainable params: 0
```

```
Epoch 1/20
18904/18904 - 70s - loss: 0.0060 - mae: 0.0527 - mape: 26.4303 - val_loss: 0.0044 - val_mae: 0.0504 - val_mape: 21.3368
Epoch 2/20
18904/18904 - 70s - loss: 0.0029 - mae: 0.0378 - mape: 19.7297 - val_loss: 0.0023 - val_mae: 0.0339 - val_mape: 16.8933
Epoch 3/20
18904/18904 - 71s - loss: 0.0025 - mae: 0.0350 - mape: 18.3935 - val_loss: 0.0022 - val_mae: 0.0338 - val_mape: 16.4769
Epoch 4/20
18904/18904 - 69s - loss: 0.0024 - mae: 0.0337 - mape: 17.7326 - val_loss: 0.0026 - val_mae: 0.0365 - val_mape: 14.7955
Epoch 5/20
18904/18904 - 70s - loss: 0.0022 - mae: 0.0323 - mape: 17.4271 - val_loss: 0.0029 - val_mae: 0.0400 - val_mape: 15.5878
Epoch 6/20
18904/18904 - 71s - loss: 0.0021 - mae: 0.0317 - mape: 16.8221 - val_loss: 0.0028 - val_mae: 0.0401 - val_mape: 15.8053
Epoch 7/20
18904/18904 - 71s - loss: 0.0021 - mae: 0.0311 - mape: 16.7360 - val_loss: 0.0020 - val_mae: 0.0314 - val_mape: 13.9622
Epoch 8/20
18904/18904 - 69s - loss: 0.0021 - mae: 0.0310 - mape: 16.8856 - val_loss: 0.0018 - val_mae: 0.0292 - val_mape: 13.2568
Epoch 9/20
18904/18904 - 69s - loss: 0.0021 - mae: 0.0308 - mape: 16.4686 - val_loss: 0.0019 - val_mae: 0.0311 - val_mape: 13.9141
Epoch 10/20
18904/18904 - 71s - loss: 0.0020 - mae: 0.0304 - mape: 16.3769 - val_loss: 0.0019 - val_mae: 0.0304 - val_mape: 13.7852
Epoch 11/20
18904/18904 - 71s - loss: 0.0020 - mae: 0.0302 - mape: 16.4885 - val_loss: 0.0019 - val_mae: 0.0305 - val_mape: 13.2742
Epoch 12/20
18904/18904 - 69s - loss: 0.0020 - mae: 0.0301 - mape: 16.3227 - val_loss: 0.0018 - val_mae: 0.0300 - val_mape: 13.2368
Epoch 13/20
18904/18904 - 71s - loss: 0.0020 - mae: 0.0301 - mape: 16.1918 - val_loss: 0.0022 - val_mae: 0.0342 - val_mape: 14.5272
Epoch 14/20
18904/18904 - 71s - loss: 0.0016 - mae: 0.0258 - mape: 15.1422 - val_loss: 0.0015 - val_mae: 0.0259 - val_mape: 12.7123
Epoch 15/20
18904/18904 - 69s - loss: 0.0016 - mae: 0.0257 - mape: 15.1254 - val_loss: 0.0016 - val_mae: 0.0271 - val_mape: 12.8926
Epoch 16/20
18904/18904 - 70s - loss: 0.0016 - mae: 0.0257 - mape: 15.1013 - val_loss: 0.0015 - val_mae: 0.0260 - val_mape: 12.6072
Epoch 17/20
18904/18904 - 71s - loss: 0.0016 - mae: 0.0257 - mape: 15.1451 - val_loss: 0.0015 - val_mae: 0.0252 - val_mape: 12.5537
Epoch 18/20
18904/18904 - 69s - loss: 0.0016 - mae: 0.0257 - mape: 15.0979 - val_loss: 0.0016 - val_mae: 0.0265 - val_mape: 12.5819
Epoch 19/20
18904/18904 - 70s - loss: 0.0016 - mae: 0.0257 - mape: 15.0406 - val_loss: 0.0015 - val_mae: 0.0252 - val_mape: 12.5789
Epoch 20/20
18904/18904 - 71s - loss: 0.0016 - mae: 0.0252 - mape: 14.8436 - val_loss: 0.0015 - val_mae: 0.0250 - val_mape: 12.4887
```

Second, check multiple time-steps.

```
h['mlp 24l 4s'] = run_model(models['mlp 24l 4s'], ds['train 24l 4s'], ds['valid 24l 4s'], optimizer = 'adadelta', epochs = 20)
```

Model: "MLP"

Layer (type)	Output Shape	Param #
flatten_3 (Flatten)	(None, 240)	0
dense_9 (Dense)	(None, 64)	15424
dense_10 (Dense)	(None, 64)	4160
dense_11 (Dense)	(None, 4)	260
reshape_3 (Reshape)	(None, 1, 4)	0

```
Total params: 19,844
Trainable params: 19,844
Non-trainable params: 0
```

```
Epoch 1/20
18869/18869 - 70s - loss: 0.0079 - mae: 0.0616 - mape: 48.5360 - val_loss: 0.0047 - val_mae: 0.0513 - val_mape: 26.9024
Epoch 2/20
18869/18869 - 70s - loss: 0.0047 - mae: 0.0489 - mape: 39.8717 - val_loss: 0.0043 - val_mae: 0.0489 - val_mape: 28.2949
Epoch 3/20
18869/18869 - 72s - loss: 0.0042 - mae: 0.0461 - mape: 36.7081 - val_loss: 0.0038 - val_mae: 0.0442 - val_mape: 26.6455
Epoch 4/20
18869/18869 - 72s - loss: 0.0040 - mae: 0.0446 - mape: 38.5765 - val_loss: 0.0046 - val_mae: 0.0512 - val_mape: 26.2144
Epoch 5/20
18869/18869 - 70s - loss: 0.0038 - mae: 0.0438 - mape: 38.0278 - val_loss: 0.0034 - val_mae: 0.0411 - val_mape: 26.2134
Epoch 6/20
18869/18869 - 70s - loss: 0.0037 - mae: 0.0428 - mape: 39.3668 - val_loss: 0.0032 - val_mae: 0.0390 - val_mape: 23.8590
Epoch 7/20
18869/18869 - 70s - loss: 0.0037 - mae: 0.0425 - mape: 36.4837 - val_loss: 0.0033 - val_mae: 0.0401 - val_mape: 22.7384
Epoch 8/20
18869/18869 - 71s - loss: 0.0036 - mae: 0.0419 - mape: 36.1011 - val_loss: 0.0036 - val_mae: 0.0427 - val_mape: 23.8585
Epoch 9/20
18869/18869 - 71s - loss: 0.0035 - mae: 0.0416 - mape: 36.2540 - val_loss: 0.0034 - val_mae: 0.0412 - val_mape: 22.7704
Epoch 10/20
18869/18869 - 72s - loss: 0.0035 - mae: 0.0415 - mape: 34.8899 - val_loss: 0.0040 - val_mae: 0.0440 - val_mape: 24.1475
Epoch 11/20
18869/18869 - 72s - loss: 0.0035 - mae: 0.0413 - mape: 33.5526 - val_loss: 0.0036 - val_mae: 0.0440 - val_mape: 25.2787
Epoch 12/20
18869/18869 - 71s - loss: 0.0031 - mae: 0.0378 - mape: 33.7398 - val_loss: 0.0030 - val_mae: 0.0375 - val_mape: 22.4531
Epoch 13/20
18869/18869 - 72s - loss: 0.0031 - mae: 0.0377 - mape: 33.9282 - val_loss: 0.0031 - val_mae: 0.0384 - val_mape: 22.5544
Epoch 14/20
18869/18869 - 70s - loss: 0.0031 - mae: 0.0377 - mape: 34.7801 - val_loss: 0.0031 - val_mae: 0.0384 - val_mape: 22.0597
Epoch 15/20
18869/18869 - 71s - loss: 0.0031 - mae: 0.0377 - mape: 34.2808 - val_loss: 0.0031 - val_mae: 0.0386 - val_mape: 22.2223
Epoch 16/20
18869/18869 - 70s - loss: 0.0030 - mae: 0.0376 - mape: 34.3899 - val_loss: 0.0031 - val_mae: 0.0388 - val_mape: 21.9613
Epoch 17/20
18869/18869 - 70s - loss: 0.0030 - mae: 0.0376 - mape: 34.5809 - val_loss: 0.0030 - val_mae: 0.0375 - val_mape: 22.0642
Epoch 18/20
18869/18869 - 71s - loss: 0.0030 - mae: 0.0373 - mape: 33.4497 - val_loss: 0.0029 - val_mae: 0.0371 - val_mape: 21.8037
Epoch 19/20
18869/18869 - 70s - loss: 0.0030 - mae: 0.0373 - mape: 33.6596 - val_loss: 0.0029 - val_mae: 0.0369 - val_mape: 21.8714
Epoch 20/20
18869/18869 - 70s - loss: 0.0030 - mae: 0.0373 - mape: 33.5927 - val_loss: 0.0029 - val_mae: 0.0370 - val_mape: 21.8664
```

- ✓ Fully convolutional network

See [Time Series Classification from Scratch with Deep Neural Networks: A Strong Baseline](#) for a detailed description of the Fully Convolutional Network (FCN) architecture. The FCN was first described in [Time-series modeling with undecimated fully convolutional neural networks](#).

The FCN architecture is a variant of the Convolutional Neural Network (CNN). A Convolutional Neural Network (CNN) usually contains fully-connected layers or a MLP at the end of the network. The FCN does not include these final layers, so it is learning convolutional filters everywhere.

TODO Include figure comparing FCNs and CNNs

The Keras [Conv1D](#) layer is used for temporal convolution.

Next, run the learning rate finder for FCNs.

```
def build_fcn_model(name, data, n_feature_maps = 64):
    in_shape, out_shape = get_io_shapes(data)
    out_steps = out_shape[0]

    fcn = Sequential(name = name)
    fcn.add(Input(shape = in_shape))

    fcn.add(Conv1D(filters = n_feature_maps, kernel_size = 8, padding = 'same'))
    fcn.add(BatchNormalization())
    fcn.add(Activation(activation = 'relu'))

    fcn.add(Conv1D(filters = n_feature_maps, kernel_size = 5, padding = 'same'))
```

```

fcn.add(BatchNormalization())
fcn.add(Activation(activation = 'relu'))

fcn.add(Conv1D(filters = n_feature_maps, kernel_size = 3, padding = 'same'))
fcn.add(BatchNormalization())
fcn.add(Activation(activation = 'relu'))

fcn.add(GlobalAveragePooling1D())
fcn.add(Dense(out_steps))

return fcn


name = 'FCN'
models['fcn_24l_1s'] = build_fcn_model(name, ds['train_24l_1s'])
models['fcn_24l_4s'] = build_fcn_model(name, ds['train_24l_4s'])

model = models['fcn_24l_1s']
model.compile(loss = 'mse', metrics = ['mae', 'mape'])
lrf_fcn_24l_1s = LRFinder(model)
lrf_fcn_24l_1s.summarise_lr(ds['train_24l_1s'], 0.001, 1, 32, 5, 50, 25)
lrf['fcn_24l_1s'] = lrf_fcn_24l_1s

model = models['fcn_24l_4s']
model.compile(loss = 'mse', metrics = ['mae', 'mape'])
lrf_fcn_24l_4s = LRFinder(model)
lrf_fcn_24l_4s.summarise_lr(ds['train_24l_4s'], 0.001, 1, 32, 5, 50, 25)
lrf['fcn_24l_4s'] = lrf_fcn_24l_4s


```

Epoch 1/5
 18904/18904 [=====] - 25s 1ms/step - loss: 0.1970 - mae: 0.2458 - mape: 114.9170



best lr: 0.0011266746

Epoch 1/5
 18869/18869 [=====] - 25s 1ms/step - loss: 0.2256 - mae: 0.2510 - mape: 124.7014



best lr: 0.0011242866

Best learning rates from the learning rate finder are close to the accepted adam learning rate of 0.001. So, I default back to that value for FCNs.

First, check single step-ahead predictions.

```
h['fcn_24l_1s'] = run_model(models['fcn_24l_1s'], ds['train_24l_1s'], ds['valid_24l_1s'], epochs = 20)
```

Model: "FCN"

Layer (type)	Output Shape	Param #
conv1d_34 (Conv1D)	(None, 24, 64)	5184
batch_normalization_36 (Batch Normalization)	(None, 24, 64)	256
activation_30 (Activation)	(None, 24, 64)	0
conv1d_35 (Conv1D)	(None, 24, 64)	20544
batch_normalization_37 (Batch Normalization)	(None, 24, 64)	256
activation_31 (Activation)	(None, 24, 64)	0
conv1d_36 (Conv1D)	(None, 24, 64)	12352
batch_normalization_38 (Batch Normalization)	(None, 24, 64)	256
activation_32 (Activation)	(None, 24, 64)	0
global_average_pooling1d_6 (Global Average Pooling1D)	(None, 64)	0
dense_18 (Dense)	(None, 1)	65

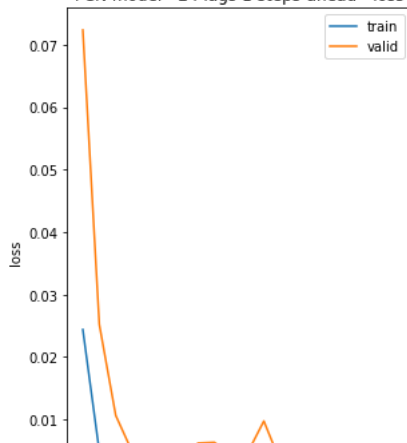
Total params: 38,913

Trainable params: 38,529

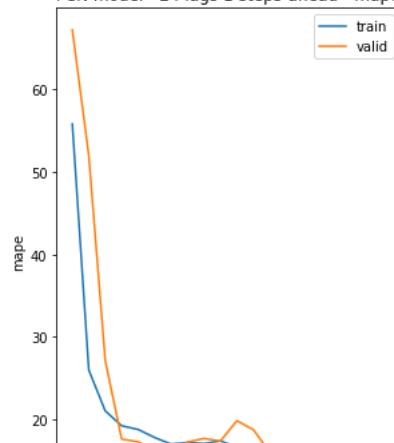
Non-trainable params: 384



Epoch 1/20
18904/18904 - 109s - loss: 0.0244 - mae: 0.1049 - mape: 55.8223 - val_loss: 0.0724 - val_mae: 0.1911 - val_mape: 67.2142
Epoch 2/20
18904/18904 - 107s - loss: 0.0049 - mae: 0.0497 - mape: 25.9807 - val_loss: 0.0253 - val_mae: 0.1245 - val_mape: 52.0181
Epoch 3/20
18904/18904 - 109s - loss: 0.0030 - mae: 0.0393 - mape: 21.0184 - val_loss: 0.0107 - val_mae: 0.0812 - val_mape: 27.1264
Epoch 4/20
18904/18904 - 109s - loss: 0.0026 - mae: 0.0357 - mape: 19.2125 - val_loss: 0.0049 - val_mae: 0.0552 - val_mape: 17.5977
Epoch 5/20
18904/18904 - 109s - loss: 0.0024 - mae: 0.0339 - mape: 18.7675 - val_loss: 0.0056 - val_mae: 0.0568 - val_mape: 17.2739
Epoch 6/20
18904/18904 - 109s - loss: 0.0022 - mae: 0.0326 - mape: 17.7924 - val_loss: 0.0046 - val_mae: 0.0528 - val_mape: 16.1029
Epoch 7/20
18904/18904 - 109s - loss: 0.0021 - mae: 0.0315 - mape: 16.9849 - val_loss: 0.0032 - val_mae: 0.0427 - val_mape: 14.8911
Epoch 8/20
18904/18904 - 109s - loss: 0.0020 - mae: 0.0310 - mape: 17.1765 - val_loss: 0.0062 - val_mae: 0.0588 - val_mape: 17.2404
Epoch 9/20
18904/18904 - 109s - loss: 0.0020 - mae: 0.0305 - mape: 17.0491 - val_loss: 0.0063 - val_mae: 0.0585 - val_mape: 17.6799
Epoch 10/20
18904/18904 - 109s - loss: 0.0020 - mae: 0.0301 - mape: 17.3789 - val_loss: 0.0042 - val_mae: 0.0458 - val_mape: 17.3576
Epoch 11/20
18904/18904 - 109s - loss: 0.0019 - mae: 0.0297 - mape: 16.5931 - val_loss: 0.0046 - val_mae: 0.0546 - val_mape: 19.8289
Epoch 12/20
18904/18904 - 109s - loss: 0.0019 - mae: 0.0296 - mape: 16.1211 - val_loss: 0.0098 - val_mae: 0.0707 - val_mape: 18.7449
Epoch 13/20
18904/18904 - 109s - loss: 0.0017 - mae: 0.0268 - mape: 15.4780 - val_loss: 0.0033 - val_mae: 0.0426 - val_mape: 15.9895
Epoch 14/20
18904/18904 - 107s - loss: 0.0017 - mae: 0.0265 - mape: 14.9915 - val_loss: 0.0025 - val_mae: 0.0370 - val_mape: 14.2585
Epoch 15/20
18904/18904 - 108s - loss: 0.0017 - mae: 0.0264 - mape: 14.9296 - val_loss: 0.0022 - val_mae: 0.0338 - val_mape: 13.7792
Epoch 16/20
18904/18904 - 108s - loss: 0.0017 - mae: 0.0264 - mape: 14.9456 - val_loss: 0.0029 - val_mae: 0.0397 - val_mape: 15.4468
Epoch 17/20
18904/18904 - 107s - loss: 0.0016 - mae: 0.0263 - mape: 15.0056 - val_loss: 0.0022 - val_mae: 0.0338 - val_mape: 13.6309
Epoch 18/20
18904/18904 - 108s - loss: 0.0016 - mae: 0.0262 - mape: 14.9100 - val_loss: 0.0027 - val_mae: 0.0378 - val_mape: 14.3438
Epoch 19/20
18904/18904 - 108s - loss: 0.0016 - mae: 0.0262 - mape: 14.7391 - val_loss: 0.0033 - val_mae: 0.0426 - val_mape: 15.4516
Epoch 20/20
18904/18904 - 107s - loss: 0.0016 - mae: 0.0261 - mape: 14.6824 - val_loss: 0.0032 - val_mae: 0.0408 - val_mape: 15.7144

FCN model - 24 lags 1 steps-ahead - loss



FCN model - 24 lags 1 steps-ahead - mape





FCN model - 24 lags 1 steps-ahead - train min loss: 0.001631 mae: 0.026096 mape: 14.682421 epoch: 20
 FCN model - 24 lags 1 steps-ahead - valid min loss: 0.002183 mae: 0.033800 mape: 13.779160 epoch: 15

Second, check multiple step-ahead predictions.

```
h['fcn_24l_4s'] = run_model(models['fcn_24l_4s'], ds['train_24l_4s'], ds['valid_24l_4s'], epochs = 20)
```

Model: "FCN"


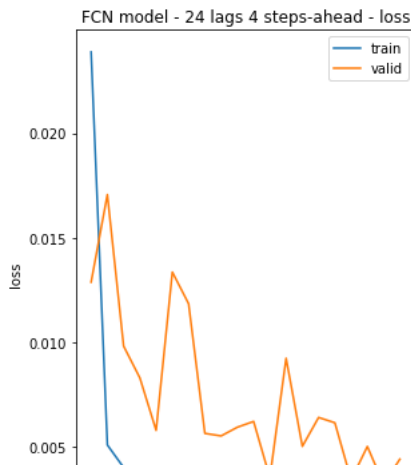
Layer (type)	Output Shape	Param #
conv1d_37 (Conv1D)	(None, 24, 64)	5184
batch_normalization_39 (Batch Normalization)	(None, 24, 64)	256
activation_33 (Activation)	(None, 24, 64)	0
conv1d_38 (Conv1D)	(None, 24, 64)	20544
batch_normalization_40 (Batch Normalization)	(None, 24, 64)	256
activation_34 (Activation)	(None, 24, 64)	0
conv1d_39 (Conv1D)	(None, 24, 64)	12352
batch_normalization_41 (Batch Normalization)	(None, 24, 64)	256
activation_35 (Activation)	(None, 24, 64)	0
global_average_pooling1d_7 (Global Average Pooling1D)	(None, 64)	0
dense_19 (Dense)	(None, 4)	260



Total params: 39,108

Trainable params: 38,724

Non-trainable params: 384

Epoch 1/20
18869/18869 - 109s - loss: 0.0239 - mae: 0.1021 - mape: 71.6267 - val_loss: 0.0129 - val_mae: 0.0876 - val_mape: 31.2475
Epoch 2/20
18869/18869 - 108s - loss: 0.0051 - mae: 0.0514 - mape: 39.6259 - val_loss: 0.0171 - val_mae: 0.1082 - val_mape: 39.5168
Epoch 3/20
18869/18869 - 108s - loss: 0.0040 - mae: 0.0452 - mape: 36.3123 - val_loss: 0.0098 - val_mae: 0.0788 - val_mape: 28.4839
Epoch 4/20
18869/18869 - 108s - loss: 0.0036 - mae: 0.0426 - mape: 33.0940 - val_loss: 0.0083 - val_mae: 0.0704 - val_mape: 23.4077
Epoch 5/20
18869/18869 - 108s - loss: 0.0034 - mae: 0.0414 - mape: 37.4728 - val_loss: 0.0058 - val_mae: 0.0586 - val_mape: 24.2517
Epoch 6/20
18869/18869 - 108s - loss: 0.0033 - mae: 0.0403 - mape: 32.2820 - val_loss: 0.0134 - val_mae: 0.0918 - val_mape: 32.4899
Epoch 7/20
18869/18869 - 108s - loss: 0.0033 - mae: 0.0400 - mape: 32.9658 - val_loss: 0.0118 - val_mae: 0.0901 - val_mape: 29.8904
Epoch 8/20
18869/18869 - 109s - loss: 0.0032 - mae: 0.0392 - mape: 33.4380 - val_loss: 0.0056 - val_mae: 0.0567 - val_mape: 22.7888
Epoch 9/20
18869/18869 - 108s - loss: 0.0031 - mae: 0.0388 - mape: 35.3697 - val_loss: 0.0055 - val_mae: 0.0578 - val_mape: 21.0027
Epoch 10/20
18869/18869 - 108s - loss: 0.0031 - mae: 0.0386 - mape: 34.3711 - val_loss: 0.0059 - val_mae: 0.0591 - val_mape: 23.2947
Epoch 11/20
18869/18869 - 108s - loss: 0.0030 - mae: 0.0382 - mape: 34.3594 - val_loss: 0.0062 - val_mae: 0.0610 - val_mape: 21.3616
Epoch 12/20
18869/18869 - 108s - loss: 0.0030 - mae: 0.0379 - mape: 36.3199 - val_loss: 0.0035 - val_mae: 0.0430 - val_mape: 20.1916
Epoch 13/20
18869/18869 - 108s - loss: 0.0029 - mae: 0.0375 - mape: 35.4463 - val_loss: 0.0092 - val_mae: 0.0798 - val_mape: 29.7993
Epoch 14/20
18869/18869 - 108s - loss: 0.0029 - mae: 0.0374 - mape: 31.0279 - val_loss: 0.0050 - val_mae: 0.0547 - val_mape: 20.1532
Epoch 15/20
18869/18869 - 108s - loss: 0.0029 - mae: 0.0370 - mape: 33.5142 - val_loss: 0.0064 - val_mae: 0.0633 - val_mape: 26.1554
Epoch 16/20
18869/18869 - 108s - loss: 0.0029 - mae: 0.0369 - mape: 33.1496 - val_loss: 0.0062 - val_mae: 0.0615 - val_mape: 21.0425
Epoch 17/20
18869/18869 - 108s - loss: 0.0028 - mae: 0.0367 - mape: 33.2664 - val_loss: 0.0034 - val_mae: 0.0431 - val_mape: 19.5610
Epoch 18/20
18869/18869 - 108s - loss: 0.0028 - mae: 0.0366 - mape: 31.8402 - val_loss: 0.0050 - val_mae: 0.0530 - val_mape: 20.5192
Epoch 19/20
18869/18869 - 108s - loss: 0.0028 - mae: 0.0364 - mape: 32.1231 - val_loss: 0.0032 - val_mae: 0.0406 - val_mape: 19.7312
Epoch 20/20
18869/18869 - 108s - loss: 0.0028 - mae: 0.0361 - mape: 34.1318 - val_loss: 0.0044 - val_mae: 0.0495 - val_mape: 20.6855





FCN model - 24 lags 4 steps-ahead - train min loss: 0.002759 mae: 0.036098 mape: 34.131786 epoch: 20
 FCN model - 24 lags 4 steps-ahead - valid min loss: 0.003212 mae: 0.040593 mape: 19.731224 epoch: 19

✓ Residual network

Residual networks, or ResNets, were originally proposed in [Deep Residual Learning for Image Recognition](#).

Residual neural networks use "identity shortcut connections" to skip over some layers. Typical ResNet models are implemented with blocks of layers that contain nonlinearities (ReLU) and batch normalization. Skipping over layers may avoid the problem of vanishing gradients, by reusing activations from a previous layer until the adjacent layer learns its weights. This should allow training networks with more layers.

TODO Include basic ResNet diagram

Again, the Keras [Conv1D](#) layer is used for temporal convolution.

Next, run the learning rate finder.

```
def build_resnet_model(name, data, n_feature_maps = 64):
    in_shape, out_shape = get_io_shapes(data)
    out_steps = out_shape[0]

    input_layer = keras.layers.Input(in_shape)

    # BLOCK 1
    conv_x = keras.layers.Conv1D(filters = n_feature_maps, kernel_size = 8, padding = 'same')(input_layer)
    conv_x = keras.layers.BatchNormalization()(conv_x)
    conv_x = keras.layers.Activation('relu')(conv_x)

    conv_y = keras.layers.Conv1D(filters = n_feature_maps, kernel_size = 5, padding = 'same')(conv_x)
    conv_y = keras.layers.BatchNormalization()(conv_y)
    conv_y = keras.layers.Activation('relu')(conv_y)

    conv_z = keras.layers.Conv1D(filters = n_feature_maps, kernel_size = 3, padding='same')(conv_y)
    conv_z = keras.layers.BatchNormalization()(conv_z)

    # expand channels for the sum
    shortcut_y = keras.layers.Conv1D(filters = n_feature_maps, kernel_size = 1, padding = 'same')(input_layer)
    shortcut_y = keras.layers.BatchNormalization()(shortcut_y)

    output_block_1 = keras.layers.add([shortcut_y, conv_z])
    output_block_1 = keras.layers.Activation('relu')(output_block_1)

    # BLOCK 2
    conv_x = keras.layers.Conv1D(filters = n_feature_maps * 2, kernel_size = 8, padding = 'same')(output_block_1)
    conv_x = keras.layers.BatchNormalization()(conv_x)
    conv_x = keras.layers.Activation('relu')(conv_x)

    conv_y = keras.layers.Conv1D(filters = n_feature_maps * 2, kernel_size = 5, padding = 'same')(conv_x)
    conv_y = keras.layers.BatchNormalization()(conv_y)
    conv_y = keras.layers.Activation('relu')(conv_y)

    conv_z = keras.layers.Conv1D(filters = n_feature_maps * 2, kernel_size = 3, padding = 'same')(conv_y)
    conv_z = keras.layers.BatchNormalization()(conv_z)

    # expand channels for the sum
    shortcut_y = keras.layers.Conv1D(filters = n_feature_maps * 2, kernel_size = 1, padding = 'same')(output_block_1)
    shortcut_y = keras.layers.BatchNormalization()(shortcut_y)

    output_block_2 = keras.layers.add([shortcut_y, conv_z])
    output_block_2 = keras.layers.Activation('relu')(output_block_2)

    # BLOCK 3
    conv_x = keras.layers.Conv1D(filters = n_feature_maps * 2, kernel_size = 8, padding = 'same')(output_block_2)
    conv_x = keras.layers.BatchNormalization()(conv_x)
    conv_x = keras.layers.Activation('relu')(conv_x)

    conv_y = keras.layers.Conv1D(filters = n_feature_maps * 2, kernel_size = 5, padding = 'same')(conv_x)
    conv_y = keras.layers.BatchNormalization()(conv_y)
    conv_y = keras.layers.Activation('relu')(conv_y)

    conv_z = keras.layers.Conv1D(filters = n_feature_maps * 2, kernel_size = 3, padding = 'same')(conv_y)
    conv_z = keras.layers.BatchNormalization()(conv_z)
```

```

# no need to expand channels because they are equal
shortcut_y = keras.layers.BatchNormalization()(output_block_2)

output_block_3 = keras.layers.add([shortcut_y, conv_z])
output_block_3 = keras.layers.Activation('relu')(output_block_3)

# FINAL
gap_layer = keras.layers.GlobalAveragePooling1D()(output_block_3)
output_layer = keras.layers.Dense(out_steps)(gap_layer)
resnet = keras.models.Model(name = name, inputs = input_layer, outputs = output_layer)

return resnet


name = 'ResNet'
models['resnet_24l_1s'] = build_resnet_model(name, ds['train_24l_1s'])
models['resnet_24l_4s'] = build_resnet_model(name, ds['train_24l_4s'])

model = models['resnet_24l_1s']
model.compile(loss = 'mse', metrics = ['mae', 'mape'])
lrf_resnet_24l_1s = LRFinder(model)
lrf_resnet_24l_1s.summarise_lr(ds['train_24l_1s'], 0.001, 10, 32, 5, 100, 50)
lrf['resnet_24l_1s'] = lrf_resnet_24l_1s

model = models['resnet_24l_4s']
model.compile(loss = 'mse', metrics = ['mae', 'mape'])
lrf_resnet_24l_4s = LRFinder(model)
lrf_resnet_24l_4s.summarise_lr(ds['train_24l_4s'], 0.001, 10, 32, 5, 100, 50)
lrf['resnet_24l_4s'] = lrf_resnet_24l_4s


```

Epoch 1/5
 18904/18904 [=====] - 75s 2ms/step - loss: 0.5768 - mae: 0.3080 - mape: 143.1178



best lr: 0.0013744514

Epoch 1/5
 18869/18869 [=====] - 40s 2ms/step - loss: 0.3239 - mae: 0.2483 - mape: 124.0002



best lr: 0.0013752629

Best learning rates from the learning rate finder are close to the accepted learning rate of 0.001. So, I default back to that value for resnets.

First, check single step-ahead predictions.

```
h['resnet_24l_1s'] = run_model(models['resnet_24l_1s'], ds['train_24l_1s'], ds['valid_24l_1s'], epochs = 20)
```

Model: "ResNet"

Layer (type)	Output Shape	Param #	Connected to
input_7 (InputLayer)	[None, 24, 10]	0	
conv1d_22 (Conv1D)	(None, 24, 64)	5184	input_7[0][0]
batch_normalization_24 (BatchNormalizatio	(None, 24, 64)	256	conv1d_22[0][0]
activation_18 (Activation)	(None, 24, 64)	0	batch_normalization_24[0][0]
conv1d_23 (Conv1D)	(None, 24, 64)	20544	activation_18[0][0]
batch_normalization_25 (BatchNormalizatio	(None, 24, 64)	256	conv1d_23[0][0]
activation_19 (Activation)	(None, 24, 64)	0	batch_normalization_25[0][0]
conv1d_25 (Conv1D)	(None, 24, 64)	704	input_7[0][0]
conv1d_24 (Conv1D)	(None, 24, 64)	12352	activation_19[0][0]
batch_normalization_27 (BatchNormalizatio	(None, 24, 64)	256	conv1d_25[0][0]
batch_normalization_26 (BatchNormalizatio	(None, 24, 64)	256	conv1d_24[0][0]
add_6 (Add)	(None, 24, 64)	0	batch_normalization_27[0][0] batch_normalization_26[0][0]
activation_20 (Activation)	(None, 24, 64)	0	add_6[0][0]
conv1d_26 (Conv1D)	(None, 24, 128)	65664	activation_20[0][0]
batch_normalization_28 (BatchNormalizatio	(None, 24, 128)	512	conv1d_26[0][0]
activation_21 (Activation)	(None, 24, 128)	0	batch_normalization_28[0][0]
conv1d_27 (Conv1D)	(None, 24, 128)	82048	activation_21[0][0]
batch_normalization_29 (BatchNormalizatio	(None, 24, 128)	512	conv1d_27[0][0]
activation_22 (Activation)	(None, 24, 128)	0	batch_normalization_29[0][0]
conv1d_29 (Conv1D)	(None, 24, 128)	8320	activation_22[0][0]
conv1d_28 (Conv1D)	(None, 24, 128)	49280	activation_22[0][0]
batch_normalization_31 (BatchNormalizatio	(None, 24, 128)	512	conv1d_29[0][0]
batch_normalization_30 (BatchNormalizatio	(None, 24, 128)	512	conv1d_28[0][0]
add_7 (Add)	(None, 24, 128)	0	batch_normalization_31[0][0] batch_normalization_30[0][0]
activation_23 (Activation)	(None, 24, 128)	0	add_7[0][0]
conv1d_30 (Conv1D)	(None, 24, 128)	131200	activation_23[0][0]

Second, check multiple step-ahead predictions.

```
activation_24 (Activation)      (None, 24, 128)      0      batch_normalization_32[0][0]
```

```
h['resnet_24l_4s'] = run_model(models['resnet_24l_4s'], ds['train_24l_4s'], ds['valid_24l_4s'], epochs = 20)
```

Model: "ResNet"

Layer (type)	Output Shape	Param #	Connected to
input_8 (InputLayer)	[(None, 24, 10)]	0	
conv1d_33 (Conv1D)	(None, 24, 64)	5184	input_8[0][0]
batch_normalization_36 (BatchNormalizatio	(None, 24, 64)	256	conv1d_33[0][0]
activation_27 (Activation)	(None, 24, 64)	0	batch_normalization_36[0][0]
conv1d_34 (Conv1D)	(None, 24, 64)	20544	activation_27[0][0]
batch_normalization_37 (BatchNormalizatio	(None, 24, 64)	256	conv1d_34[0][0]
activation_28 (Activation)	(None, 24, 64)	0	batch_normalization_37[0][0]
conv1d_36 (Conv1D)	(None, 24, 64)	704	input_8[0][0]
conv1d_35 (Conv1D)	(None, 24, 64)	12352	activation_28[0][0]
batch_normalization_39 (BatchNormalizatio	(None, 24, 64)	256	conv1d_36[0][0]
batch_normalization_38 (BatchNormalizatio	(None, 24, 64)	256	conv1d_35[0][0]
add_9 (Add)	(None, 24, 64)	0	batch_normalization_39[0][0] batch_normalization_38[0][0]
activation_29 (Activation)	(None, 24, 64)	0	add_9[0][0]
conv1d_37 (Conv1D)	(None, 24, 128)	65664	activation_29[0][0]
batch_normalization_40 (BatchNormalizatio	(None, 24, 128)	512	conv1d_37[0][0]
activation_30 (Activation)	(None, 24, 128)	0	batch_normalization_40[0][0]
conv1d_38 (Conv1D)	(None, 24, 128)	82048	activation_30[0][0]
batch_normalization_41 (BatchNormalizatio	(None, 24, 128)	512	conv1d_38[0][0]
activation_31 (Activation)	(None, 24, 128)	0	batch_normalization_41[0][0]
conv1d_40 (Conv1D)	(None, 24, 128)	8320	activation_29[0][0]
conv1d_39 (Conv1D)	(None, 24, 128)	49280	activation_31[0][0]
batch_normalization_43 (BatchNormalizatio	(None, 24, 128)	512	conv1d_40[0][0]
batch_normalization_42 (BatchNormalizatio	(None, 24, 128)	512	conv1d_39[0][0]
add_10 (Add)	(None, 24, 128)	0	batch_normalization_43[0][0] batch_normalization_42[0][0]
activation_32 (Activation)	(None, 24, 128)	0	add_10[0][0]
conv1d_41 (Conv1D)	(None, 24, 128)	131200	activation_32[0][0]
batch_normalization_44 (BatchNormalizatio	(None, 24, 128)	512	conv1d_41[0][0]
activation_33 (Activation)	(None, 24, 128)	0	batch_normalization_44[0][0]
conv1d_42 (Conv1D)	(None, 24, 128)	82048	activation_33[0][0]
batch_normalization_45 (BatchNormalizatio	(None, 24, 128)	512	conv1d_42[0][0]
activation_34 (Activation)	(None, 24, 128)	0	batch_normalization_45[0][0]
conv1d_43 (Conv1D)	(None, 24, 128)	49280	activation_34[0][0]
batch_normalization_47 (BatchNormalizatio	(None, 24, 128)	512	activation_32[0][0]
batch_normalization_46 (BatchNormalizatio	(None, 24, 128)	512	conv1d_43[0][0]
add_11 (Add)	(None, 24, 128)	0	batch_normalization_47[0][0] batch_normalization_46[0][0]
activation_35 (Activation)	(None, 24, 128)	0	add_11[0][0]
global_average_pooling1d_3 (GlobalAverage	(None, 128)	0	activation_35[0][0]
dense_15 (Dense)	(None, 4)	516	global_average_pooling1d_3[0][0]

Total params: 512,260
Trainable params: 509,700
Non-trainable params: 2,560

Epoch 1/20

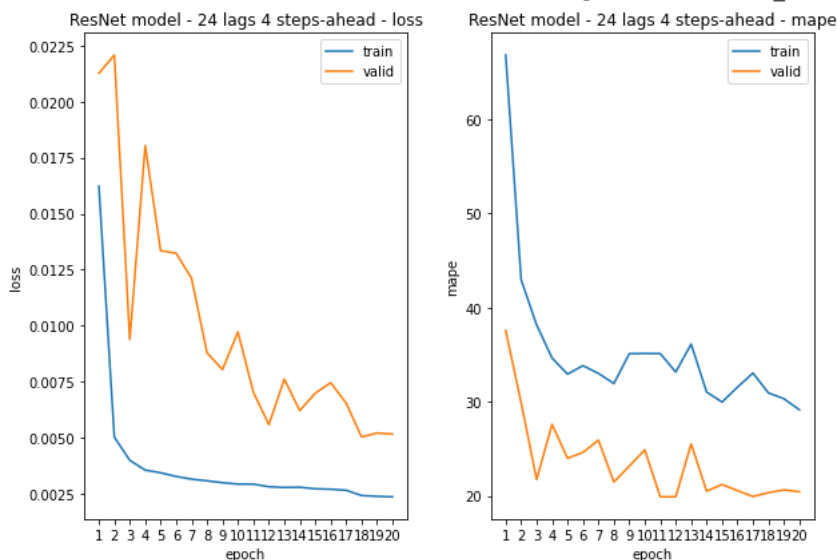
18869/18869 - 257s - loss: 0.0162 - mae: 0.0850 - mape: 66.8201 - val_loss: 0.0213 - val_mae: 0.1162 - val_mape: 37.5702

Epoch 2/20

```

18869/18869 - 256s - loss: 0.0050 - mae: 0.0507 - mape: 42.9599 - val_loss: 0.0221 - val_mae: 0.1123 - val_mape: 29.9136
Epoch 3/20
18869/18869 - 258s - loss: 0.0040 - mae: 0.0449 - mape: 38.1732 - val_loss: 0.0094 - val_mae: 0.0741 - val_mape: 21.7524
Epoch 4/20
18869/18869 - 257s - loss: 0.0036 - mae: 0.0422 - mape: 34.6486 - val_loss: 0.0180 - val_mae: 0.1072 - val_mape: 27.6115
Epoch 5/20
18869/18869 - 257s - loss: 0.0034 - mae: 0.0412 - mape: 32.9408 - val_loss: 0.0133 - val_mae: 0.0916 - val_mape: 24.0021
Epoch 6/20
18869/18869 - 258s - loss: 0.0033 - mae: 0.0400 - mape: 33.8418 - val_loss: 0.0132 - val_mae: 0.0891 - val_mape: 24.6326
Epoch 7/20
18869/18869 - 255s - loss: 0.0031 - mae: 0.0392 - mape: 33.0154 - val_loss: 0.0121 - val_mae: 0.0817 - val_mape: 25.9249
Epoch 8/20
18869/18869 - 257s - loss: 0.0031 - mae: 0.0386 - mape: 31.9466 - val_loss: 0.0088 - val_mae: 0.0711 - val_mape: 21.4855
Epoch 9/20
18869/18869 - 258s - loss: 0.0030 - mae: 0.0380 - mape: 35.1108 - val_loss: 0.0080 - val_mae: 0.0680 - val_mape: 23.1770
Epoch 10/20
18869/18869 - 254s - loss: 0.0029 - mae: 0.0375 - mape: 35.1308 - val_loss: 0.0097 - val_mae: 0.0741 - val_mape: 24.8891
Epoch 11/20
18869/18869 - 256s - loss: 0.0029 - mae: 0.0374 - mape: 35.1222 - val_loss: 0.0070 - val_mae: 0.0632 - val_mape: 19.9127
Epoch 12/20
18869/18869 - 257s - loss: 0.0028 - mae: 0.0366 - mape: 33.1597 - val_loss: 0.0056 - val_mae: 0.0561 - val_mape: 19.9193
Epoch 13/20
18869/18869 - 256s - loss: 0.0028 - mae: 0.0364 - mape: 36.1271 - val_loss: 0.0076 - val_mae: 0.0669 - val_mape: 25.5291
Epoch 14/20
18869/18869 - 254s - loss: 0.0028 - mae: 0.0364 - mape: 31.0407 - val_loss: 0.0062 - val_mae: 0.0608 - val_mape: 20.5167
Epoch 15/20
18869/18869 - 255s - loss: 0.0027 - mae: 0.0359 - mape: 29.9716 - val_loss: 0.0070 - val_mae: 0.0643 - val_mape: 21.2172
Epoch 16/20
18869/18869 - 256s - loss: 0.0027 - mae: 0.0357 - mape: 31.5652 - val_loss: 0.0075 - val_mae: 0.0673 - val_mape: 20.5711
Epoch 17/20
18869/18869 - 255s - loss: 0.0027 - mae: 0.0353 - mape: 33.0526 - val_loss: 0.0065 - val_mae: 0.0627 - val_mape: 19.9382
Epoch 18/20
18869/18869 - 257s - loss: 0.0024 - mae: 0.0331 - mape: 30.9467 - val_loss: 0.0050 - val_mae: 0.0545 - val_mape: 20.3570
Epoch 19/20
18869/18869 - 257s - loss: 0.0024 - mae: 0.0328 - mape: 30.3295 - val_loss: 0.0052 - val_mae: 0.0557 - val_mape: 20.6455
Epoch 20/20
18869/18869 - 257s - loss: 0.0024 - mae: 0.0326 - mape: 29.1513 - val_loss: 0.0052 - val_mae: 0.0553 - val_mape: 20.4484

```



```

ResNet model - 24 lags 4 steps-ahead - train min loss: 0.002362 mae: 0.032584 mape: 29.151272 epoch: 20
ResNet model - 24 lags 4 steps-ahead - valid min loss: 0.005034 mae: 0.054451 mape: 20.356956 epoch: 18

```

✓ LSTM network

Long Short Term Memory networks, or LSTMs, were originally proposed in [LONG SHORT TERM MEMORY](#). They are recurrent neural networks which have feedback connections.

LSTMs can take entire sequences of data as input and keep track of arbitrary long-term dependencies. A LSTM unit is composed of a cell and three gates. The cell remembers values over arbitrary time intervals and the input, output and forget gates regulate the flow of information into and out of the cell.

TODO Include basic LSTM diagram

Next, run the learning rate finder.

```

def build_lstm_model(name, data, n_feature_maps = 8):
    in_shape, out_shape = get_io_shapes(data)
    out_steps = out_shape[0]

```

```

lstm = Sequential(name = name)
lstm.add(Input(shape = in_shape))

# Shape [batch, time, features] => [batch, n_feature_maps]
lstm.add(LSTM(n_feature_maps, return_sequences=False))

# Shape => [batch, out_steps]
lstm.add(Dense(out_steps,
               kernel_initializer=tf.initializers.zeros()))

return lstm


name = 'LSTM'
models['lstm_24l_1s'] = build_lstm_model(name, ds['train_24l_1s'])
models['lstm_24l_4s'] = build_lstm_model(name, ds['train_24l_4s'])

model = models['lstm_24l_1s']
model.compile(loss = 'mse', metrics = ['mae', 'mape'])
lrf_lstm_24l_1s = LRFinder(model)
lrf_lstm_24l_1s.summarise_lr(ds['train_24l_1s'], 0.0001, 10, 32, 5, 100, 25)
lrf['lstm_24l_1s'] = lrf_lstm_24l_1s

model = models['lstm_24l_4s']
model.compile(loss = 'mse', metrics = ['mae', 'mape'])
lrf_lstm_24l_4s = LRFinder(model)
lrf_lstm_24l_4s.summarise_lr(ds['train_24l_4s'], 0.0001, 10, 32, 5, 100, 25)
lrf['lstm_24l_4s'] = lrf_lstm_24l_4s


```

Epoch 1/5
18904/18904 [=====] - 19s 943us/step - loss: 0.7534 - mae: 0.4657 - mape: 167.8633



best lr: 0.00022147449

Epoch 1/5
18869/18869 [=====] - 20s 966us/step - loss: 0.7260 - mae: 0.4561 - mape: 175.8322



best lr: 0.00022354048